



# **SUPERVISED LEARNING CLASSIFICATION**

# WHAT WILL WE COVER?

- **Recap of what is supervised learning**
- **Core Concepts required**
- **Classification Algorithms**
- **Regression Algorithms**

# MACHINE LEARNING: SUPERVISED VS UNSUPERVISED



**Whenever we have a defined output to predict it is a supervised ML problem**

**If we don't then we have a unsupervised learning problem**

# SUPERVISED: SHOULD I PLAY GOLF?

Input

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY
0	Rainy	Hot	High	False
1	Rainy	Hot	High	True
2	Overcast	Hot	High	False
3	Sunny	Mild	High	False
4	Sunny	Cool	Normal	False
5	Sunny	Cool	Normal	True
6	Overcast	Cool	Normal	True
7	Rainy	Mild	High	False
8	Rainy	Cool	Normal	False
9	Sunny	Mild	Normal	False
10	Rainy	Mild	Normal	True
11	Overcast	Mild	High	True
12	Overcast	Hot	Normal	False
13	Sunny	Mild	High	True

The objective is to predict if based on the weather conditions of a particular day, we should go play Golf?

# SUPERVISED: SHOULD I PLAY GOLF?

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY	PLAY GOLF
0	Rainy	Hot	High	False	No
1	Rainy	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Sunny	Mild	High	False	Yes
4	Sunny	Cool	Normal	False	Yes
5	Sunny	Cool	Normal	True	No
6	Overcast	Cool	Normal	True	Yes
7	Rainy	Mild	High	False	No
8	Rainy	Cool	Normal	False	Yes
9	Sunny	Mild	Normal	False	Yes
10	Rainy	Mild	Normal	True	Yes
11	Overcast	Mild	High	True	Yes
12	Overcast	Hot	Normal	False	Yes
13	Sunny	Mild	High	True	No

But how does my algorithm learn?

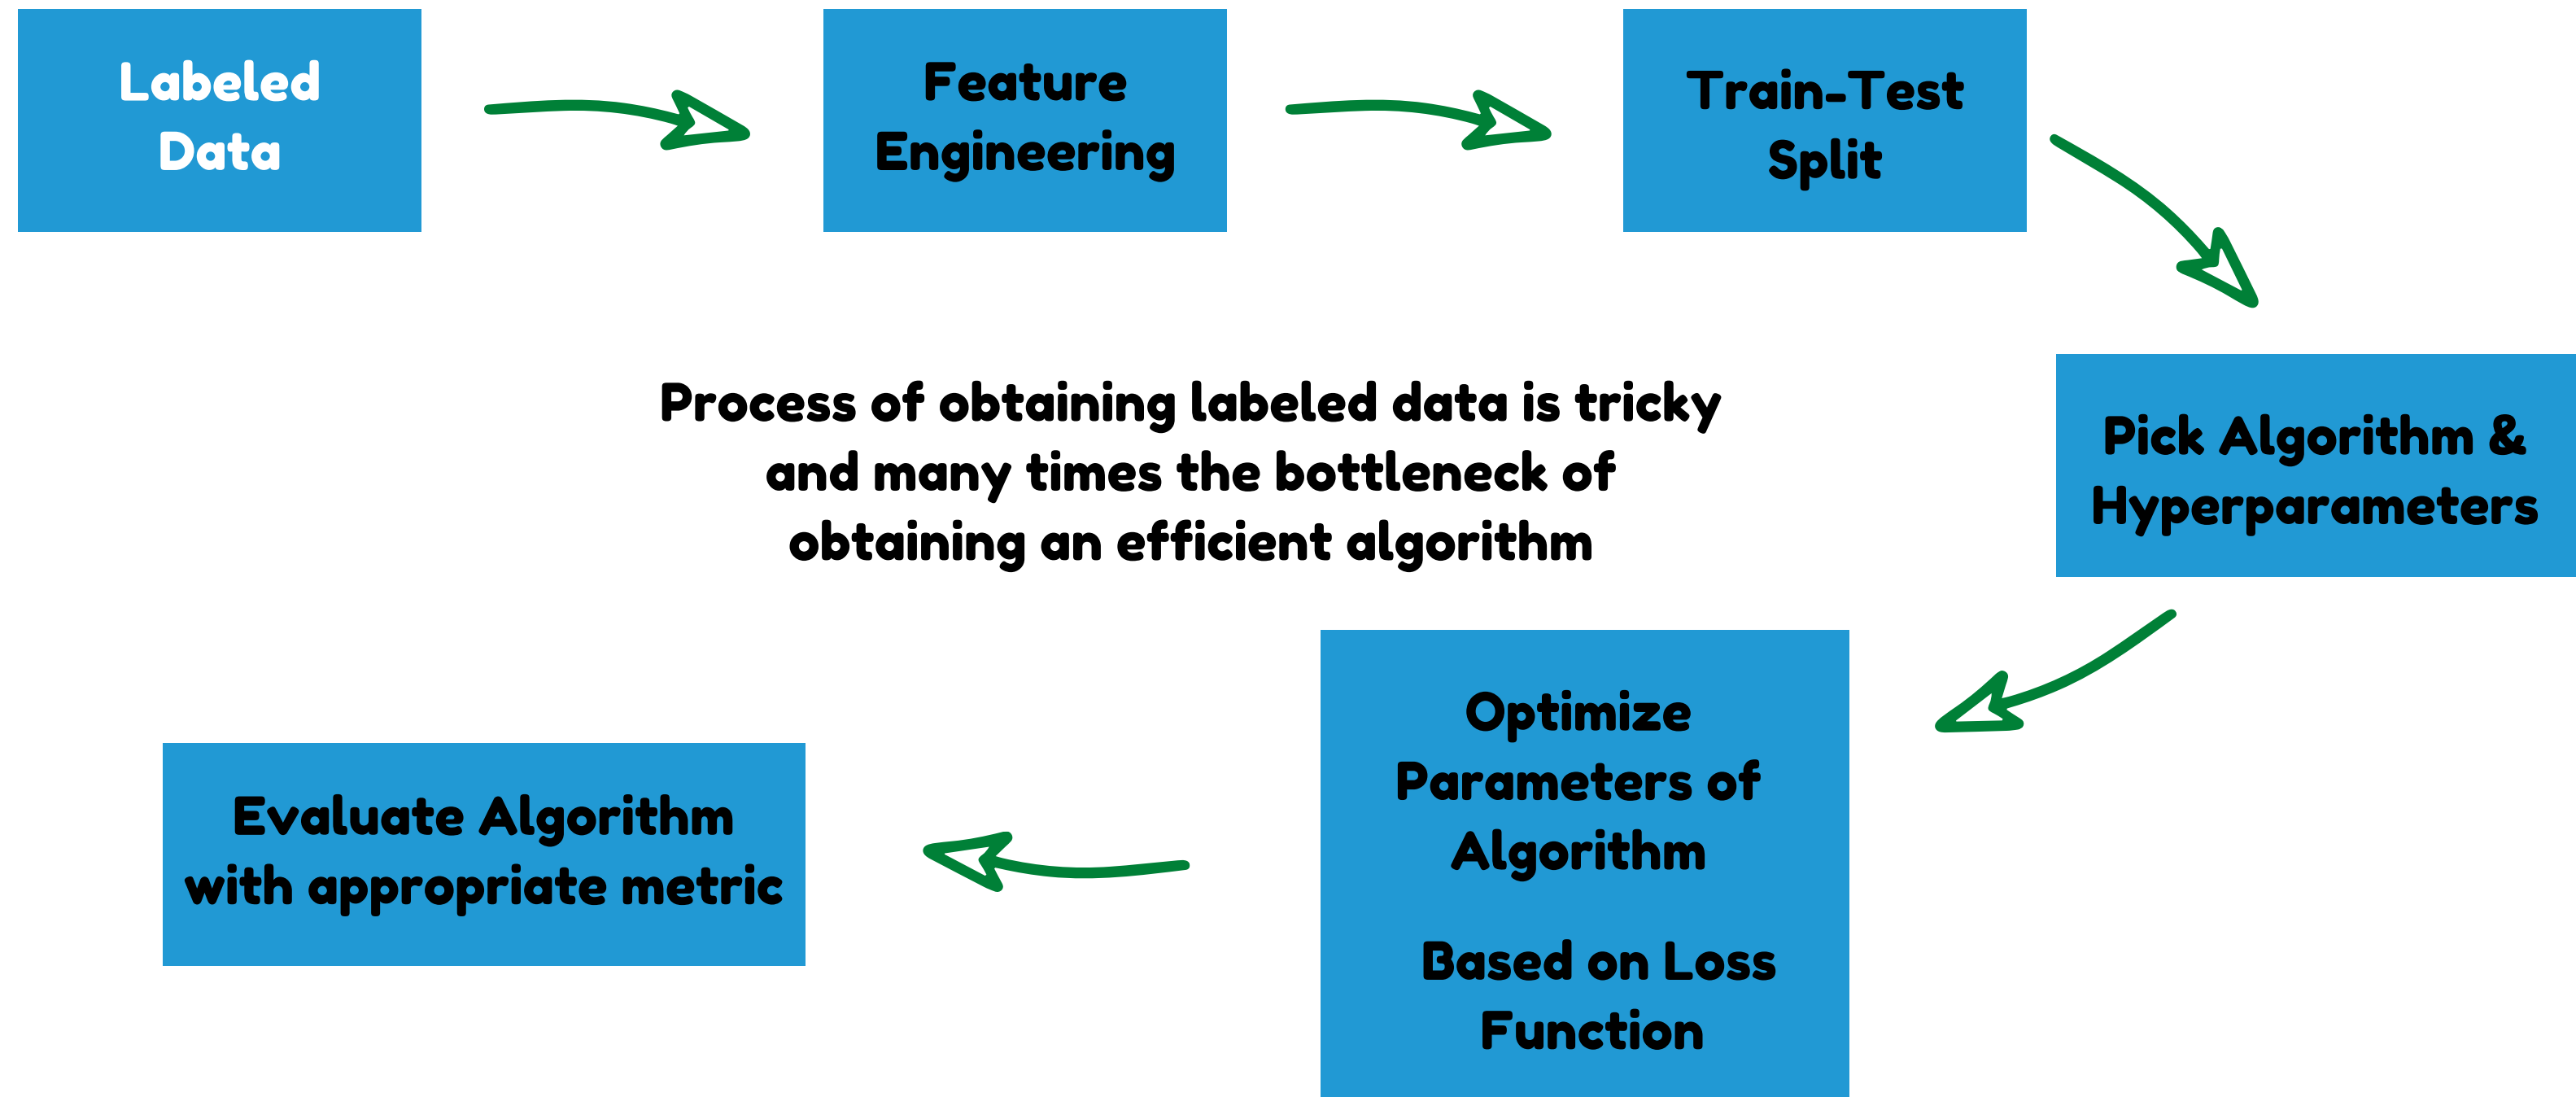
Based on past datapoints!!!

This is the "Supervised" component

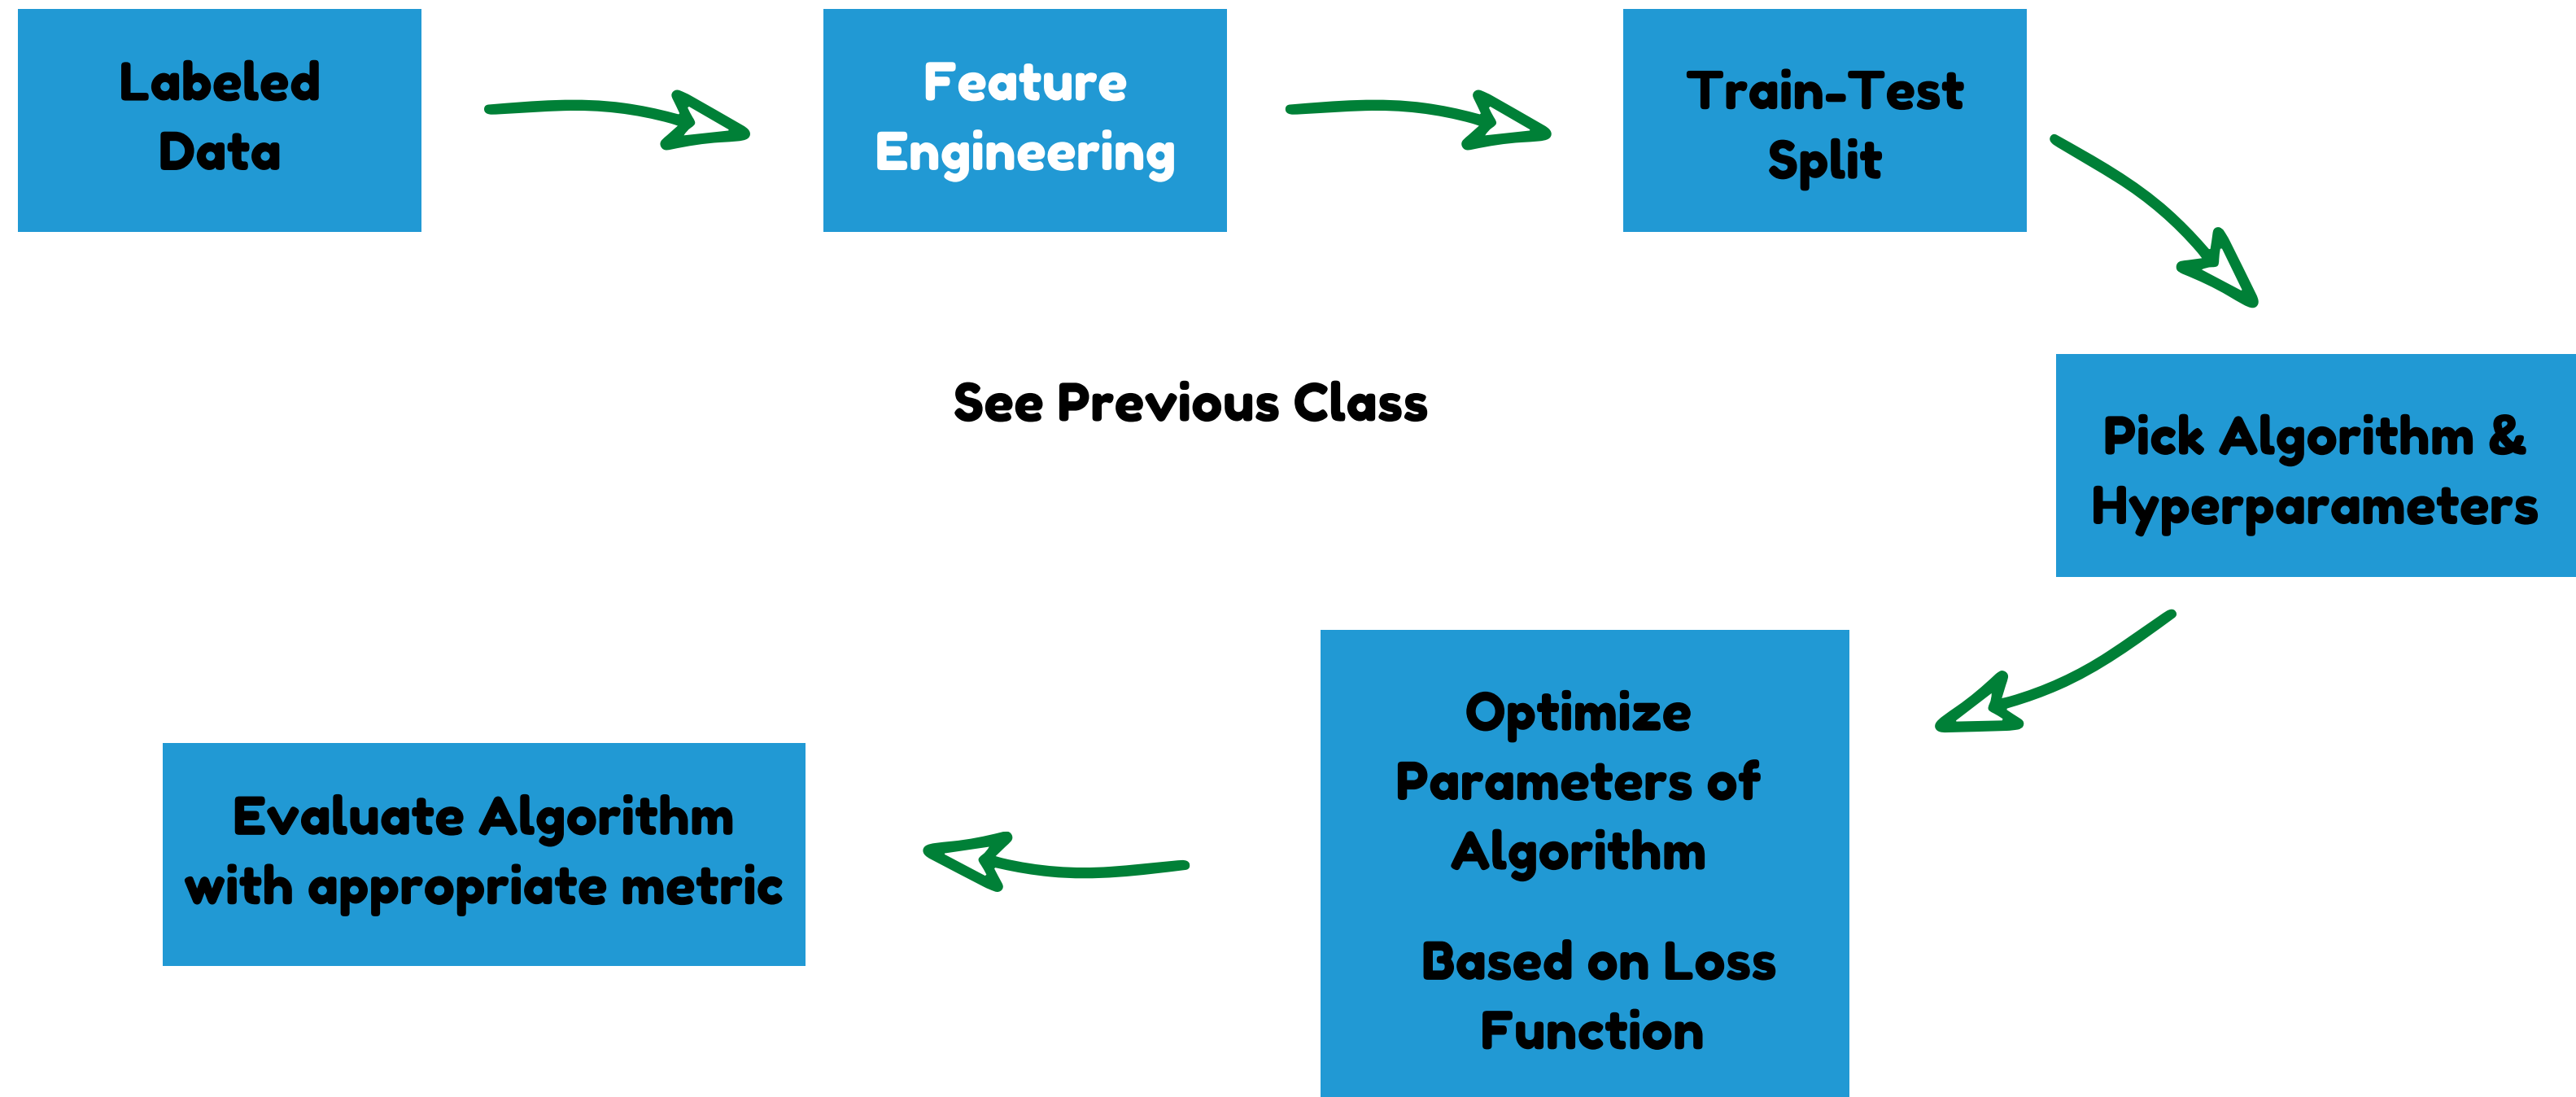
These are called the labels

Supervised Learning is performed on labelled data

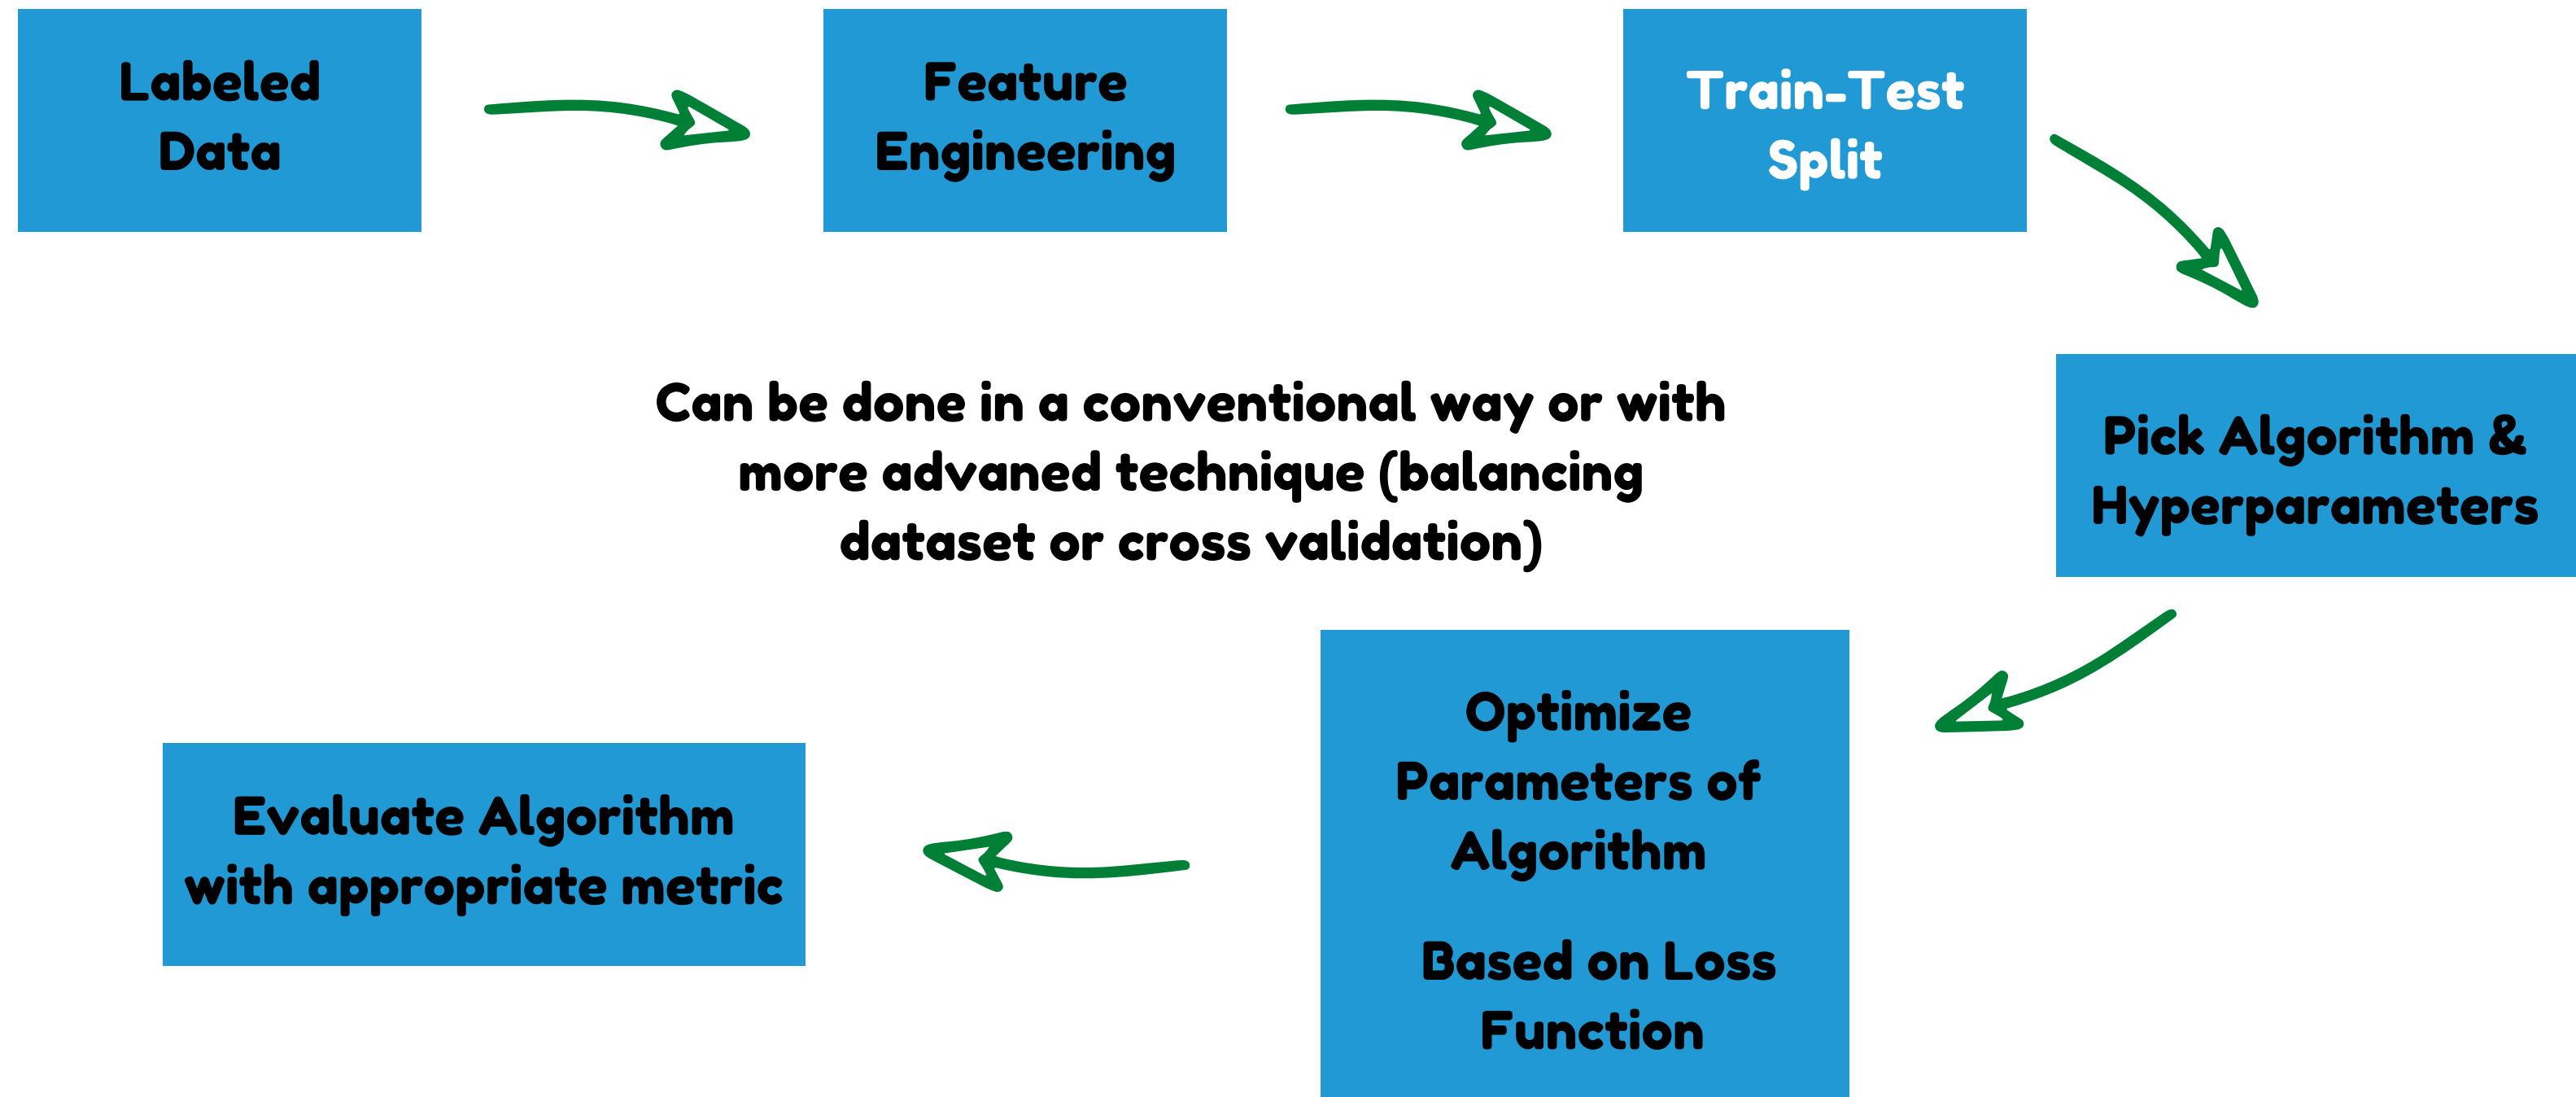
# THE MACHINE LEARNING PROCESS



# THE MACHINE LEARNING PROCESS

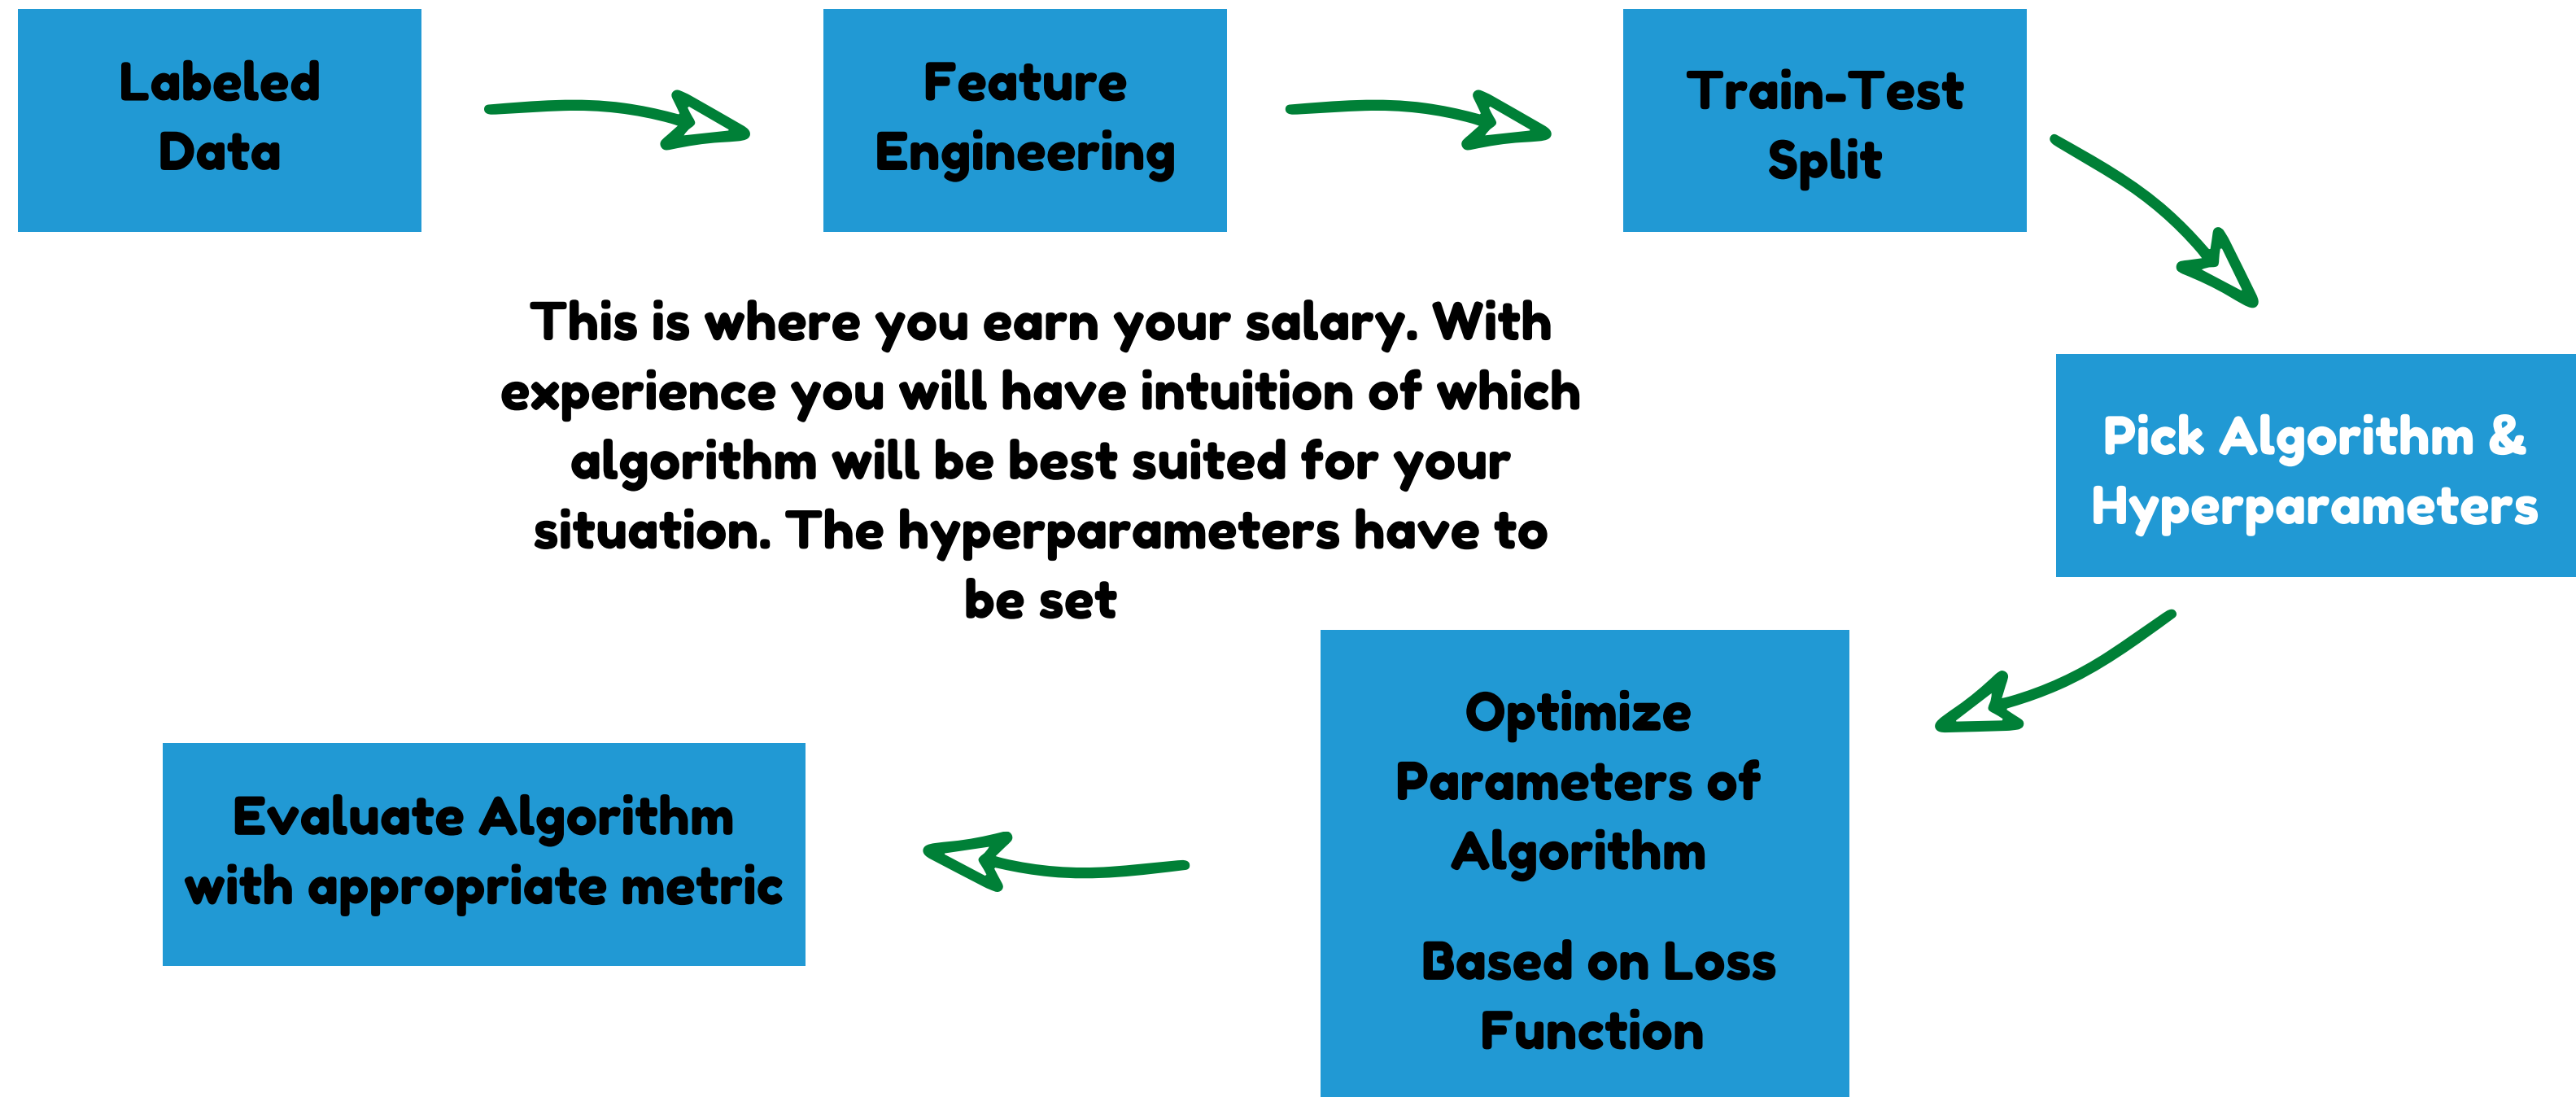


# THE MACHINE LEARNING PROCESS

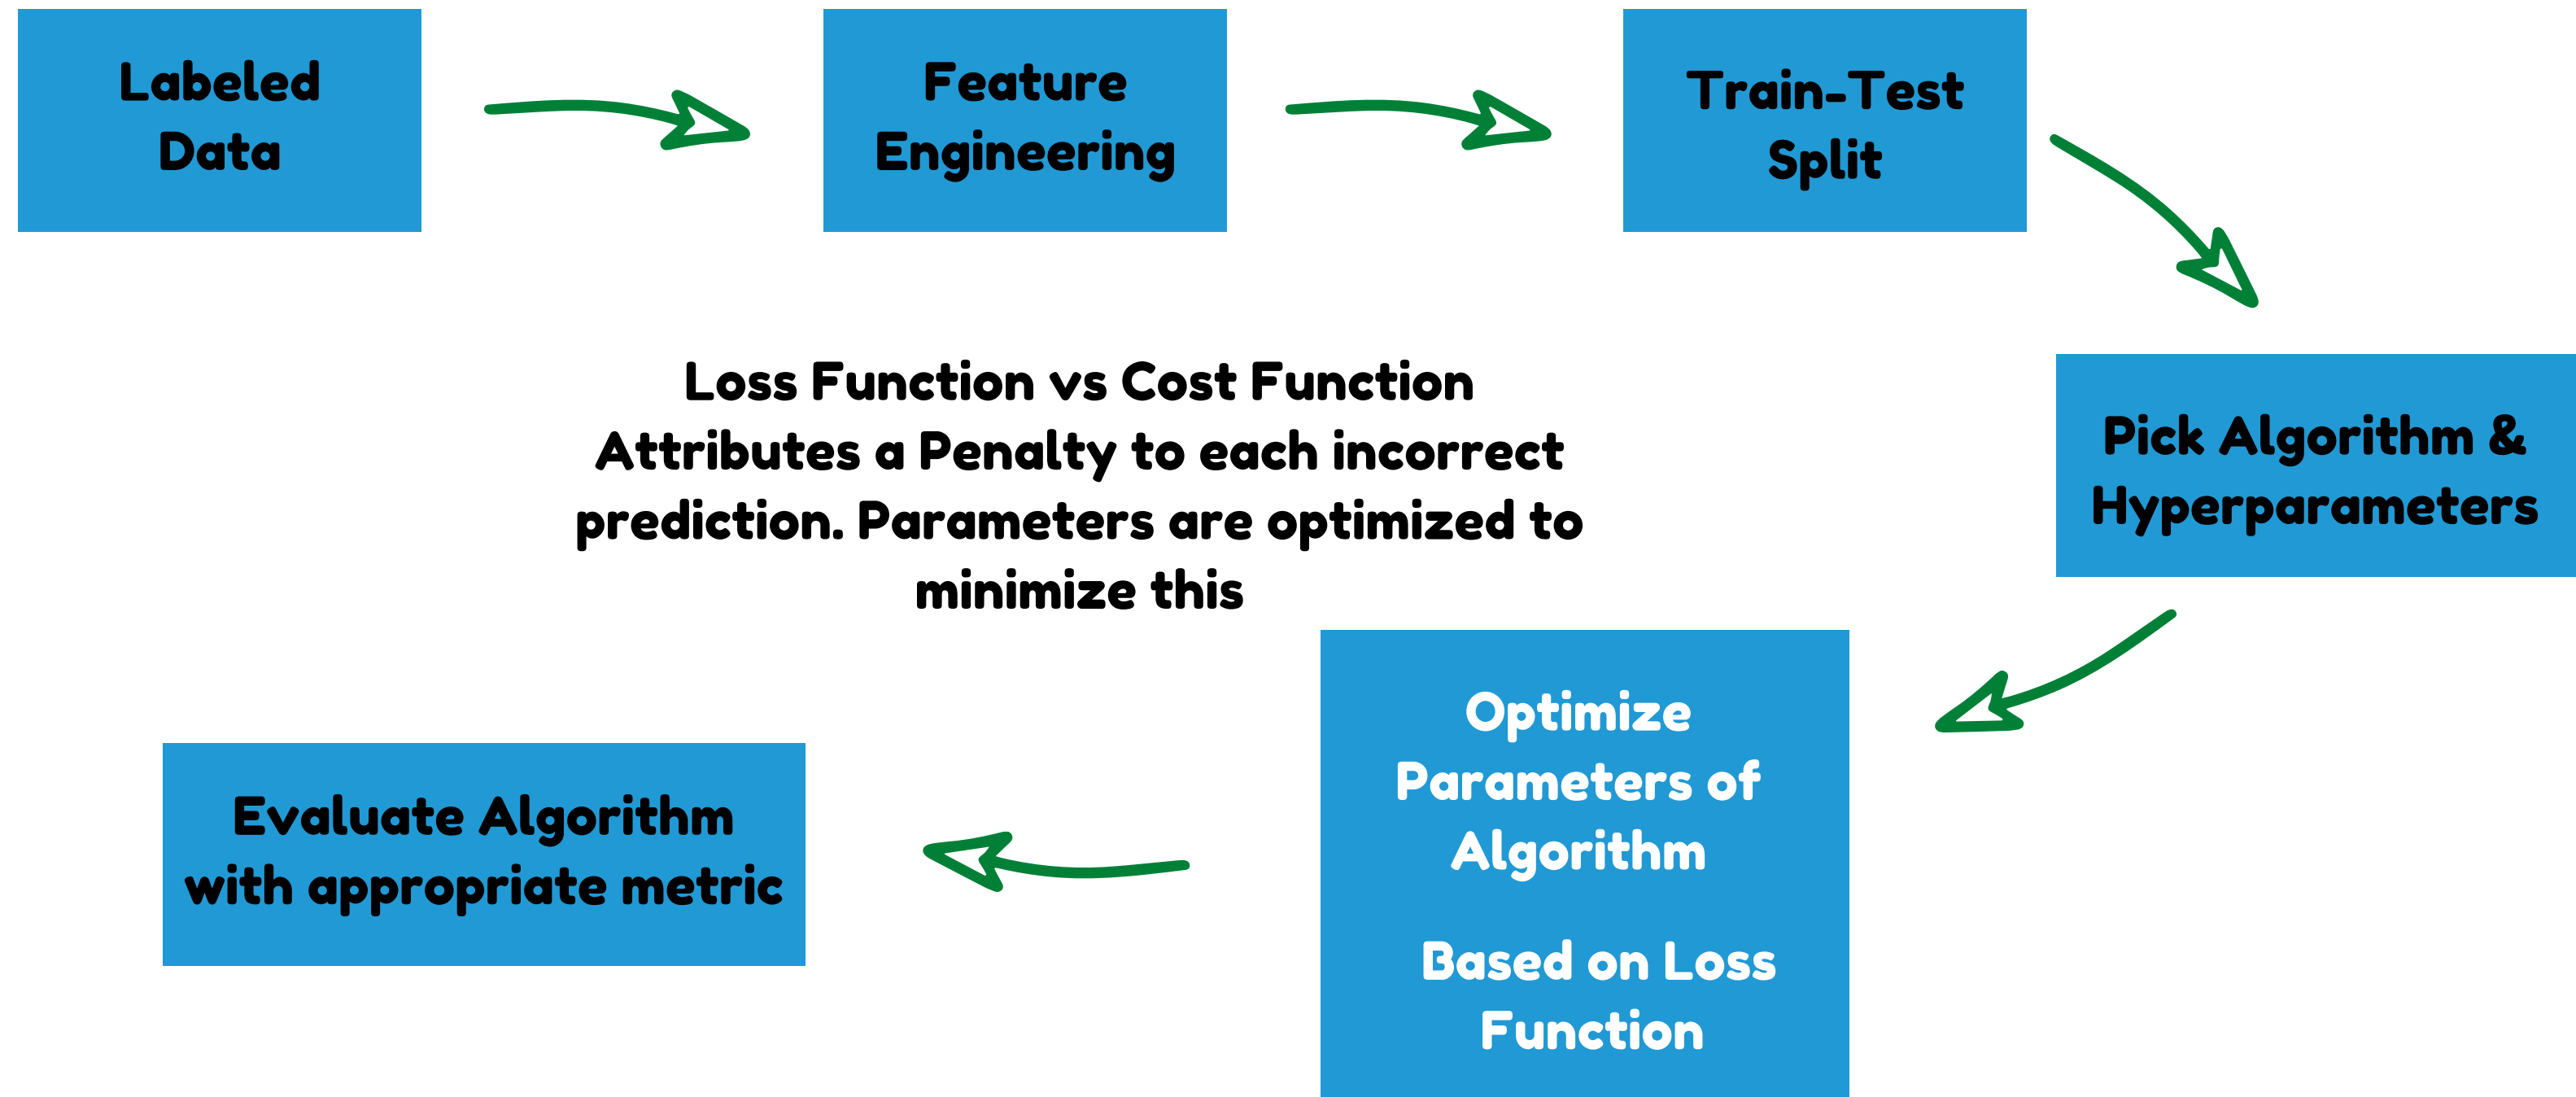




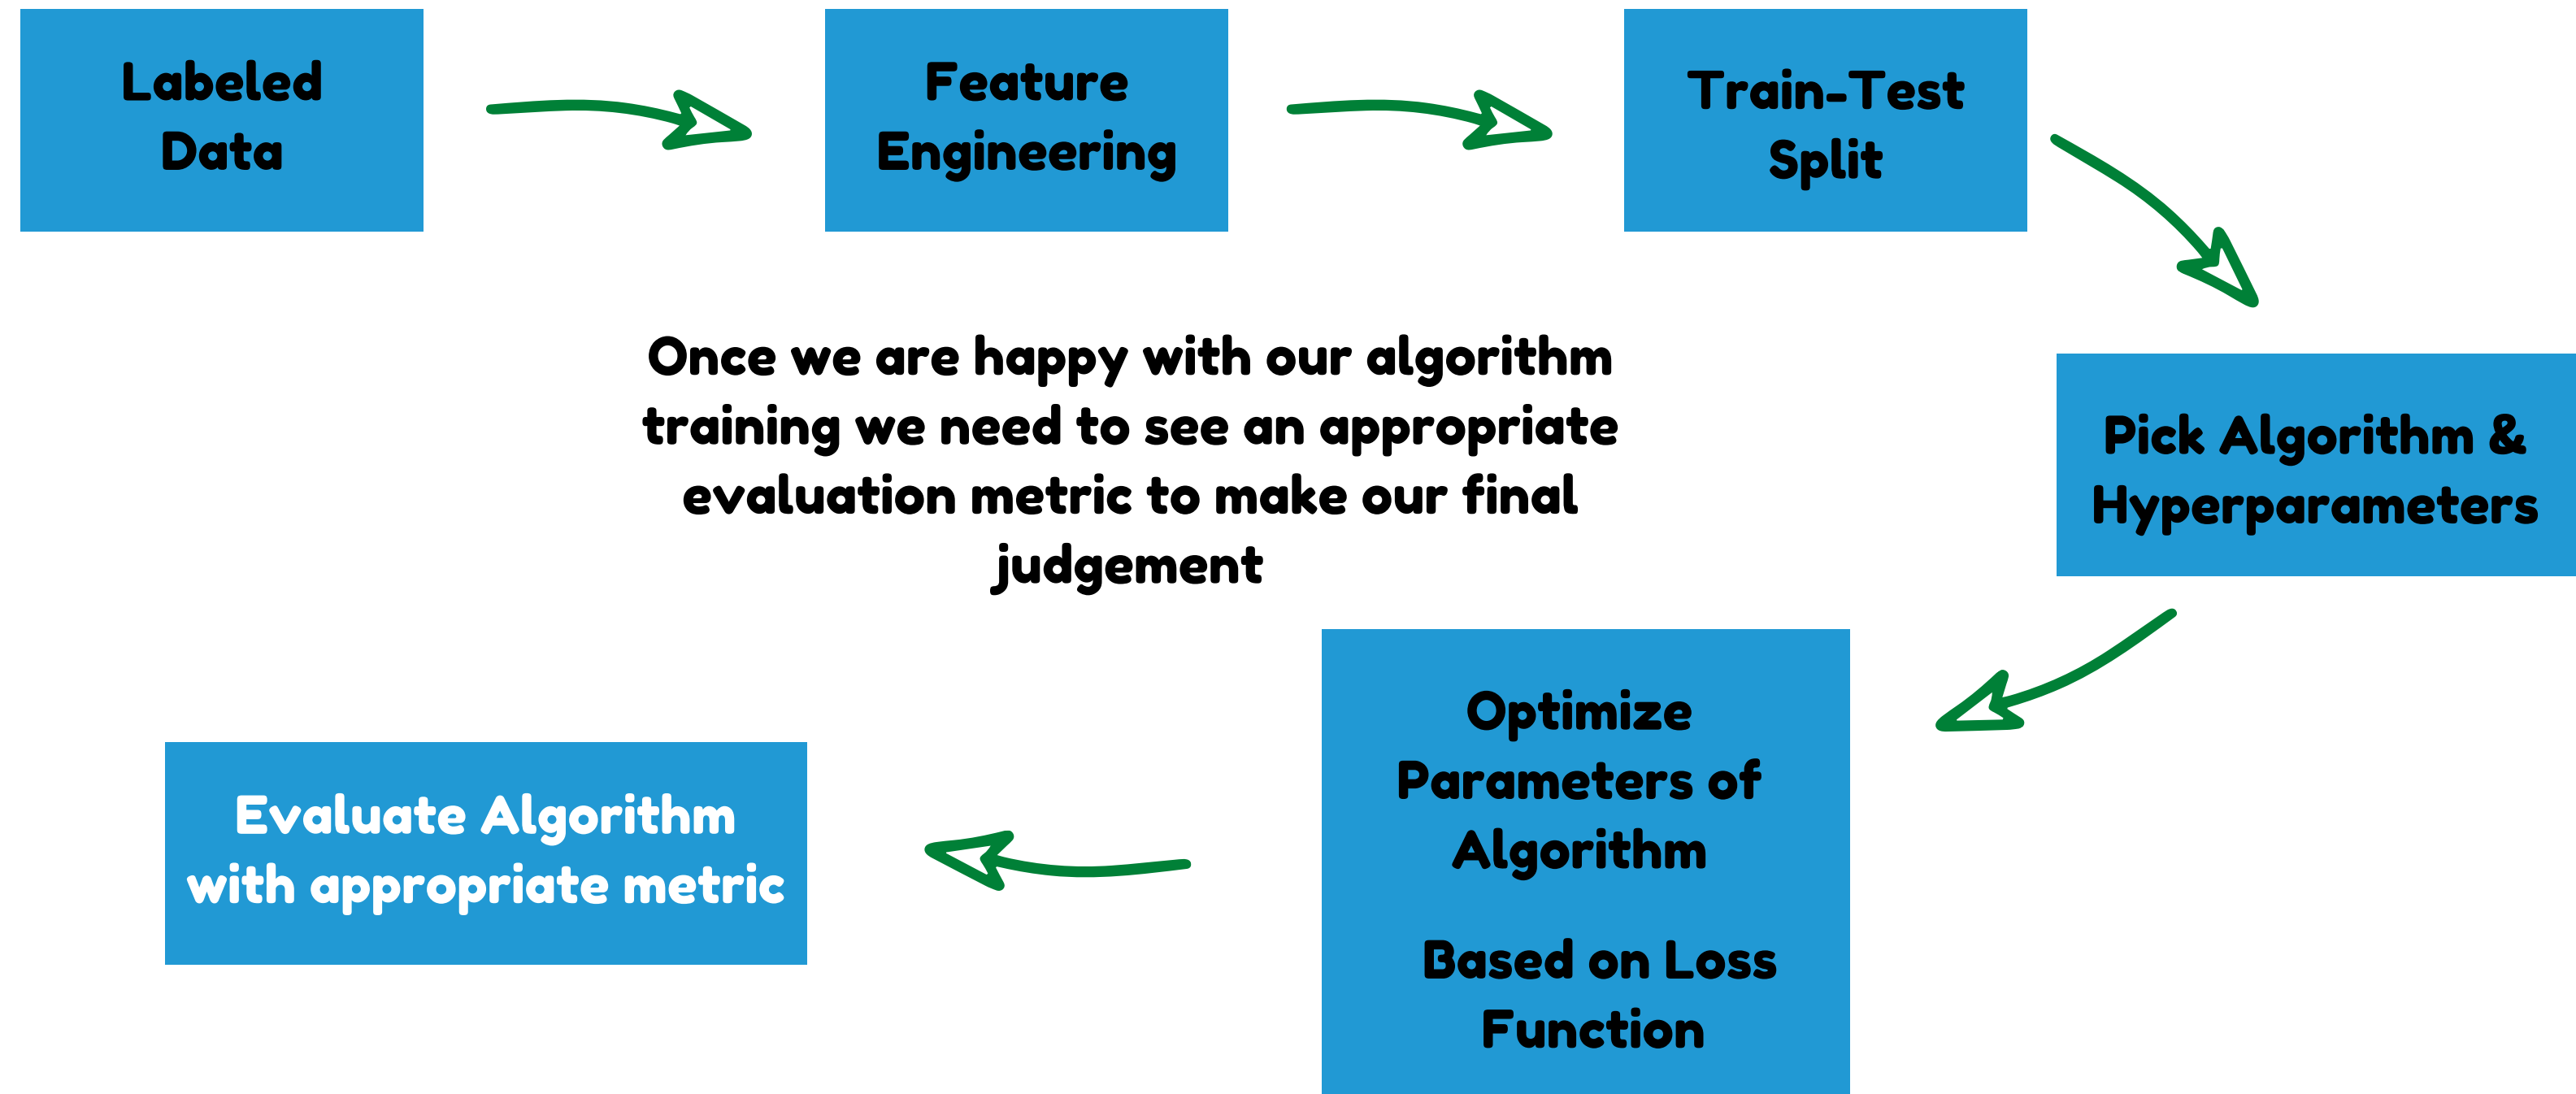
# THE MACHINE LEARNING PROCESS



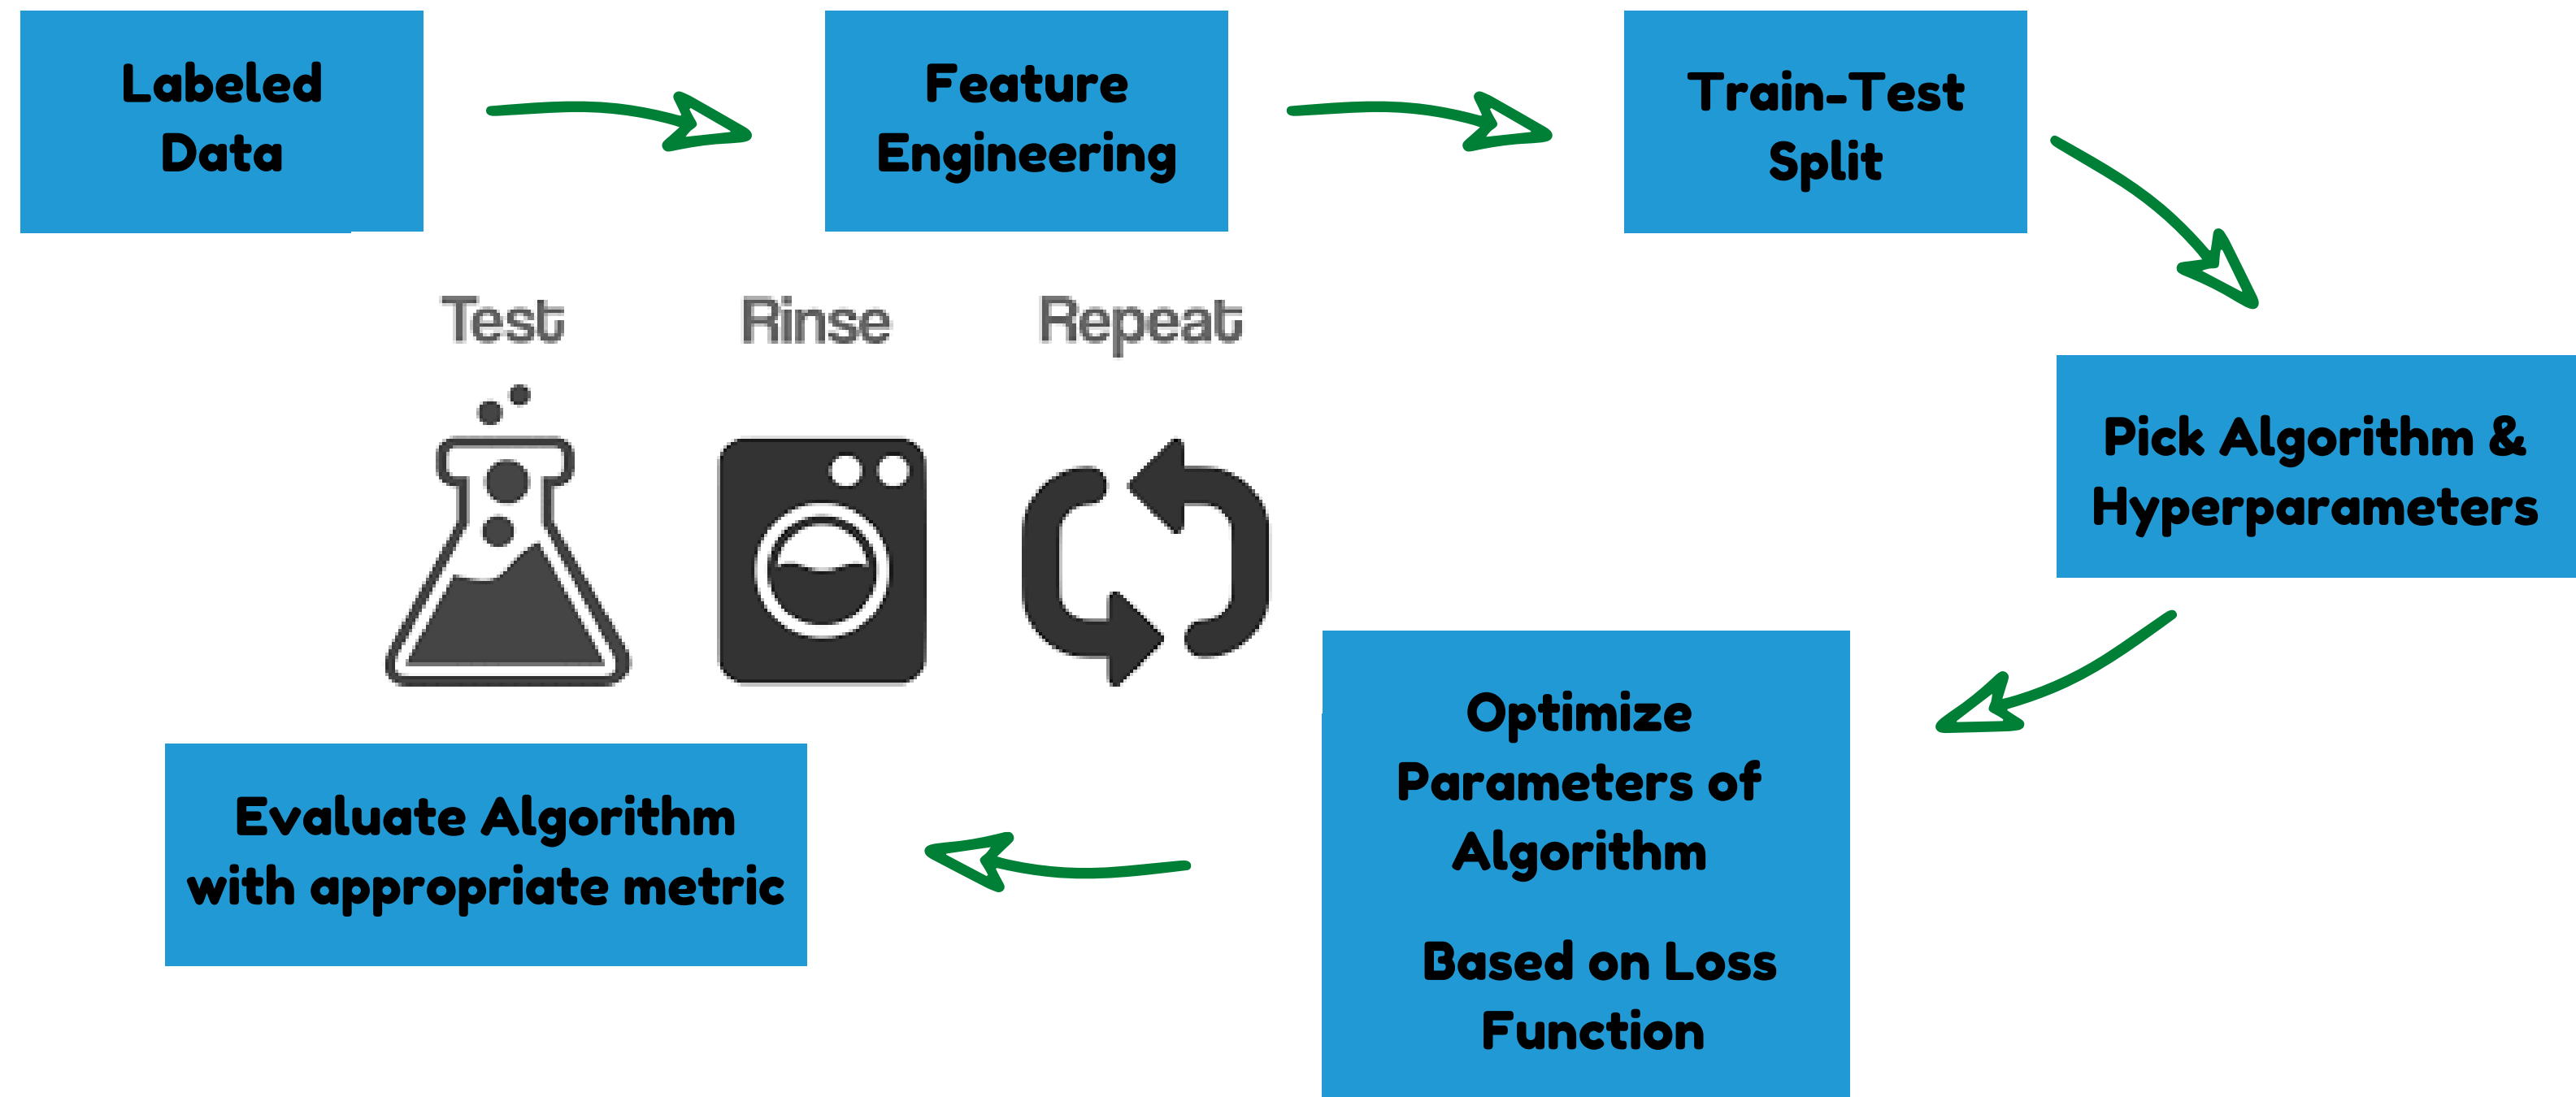
# THE MACHINE LEARNING PROCESS



# THE MACHINE LEARNING PROCESS




# THE MACHINE LEARNING PROCESS



# CLASSIFICATION ALGORITHMS

- **KNN (Regression Friendly) (Multi Class)**
- **Logistic Regression**
- **Decision Trees (Regression Friendly) (Multi Class)**
- **Naive-Bayes Classifier (Multi Class)**
- **Support Vector Machine (SVM)**

# CLASSIFICATION ALGORITHMS




- This is the biggest danger when creating machine learning algorithms

**Underfitting: model is too simplistic**

**Overfitting: model adapts too much to the training data and does not generalize properly**

# CLASSIFICATION ALGORITHMS



- **Model Generalization is how well will the trained model perform on new unseen data?**
- **This is why we perform the train-test split!!!**
- **This way we give the algorithm some data to train, and then a smaller sample to test how does the model behave on unseen data**



# K NEAREST NEIGHBOURS

## K- Nearest Neighbours (KNN)

The Copy-Cat of the Neighbours

Hyperparameters:


- K - # of neighbours
- Method of weight of neighbours





# K NEAREST NEIGHBOURS


**Example of Overfitting and training vs test tradeoff**



# KNN - SUMMARY



**No assumptions about data – useful, for example, for nonlinear data**  
**Simple algorithm – to explain and understand/interpret**



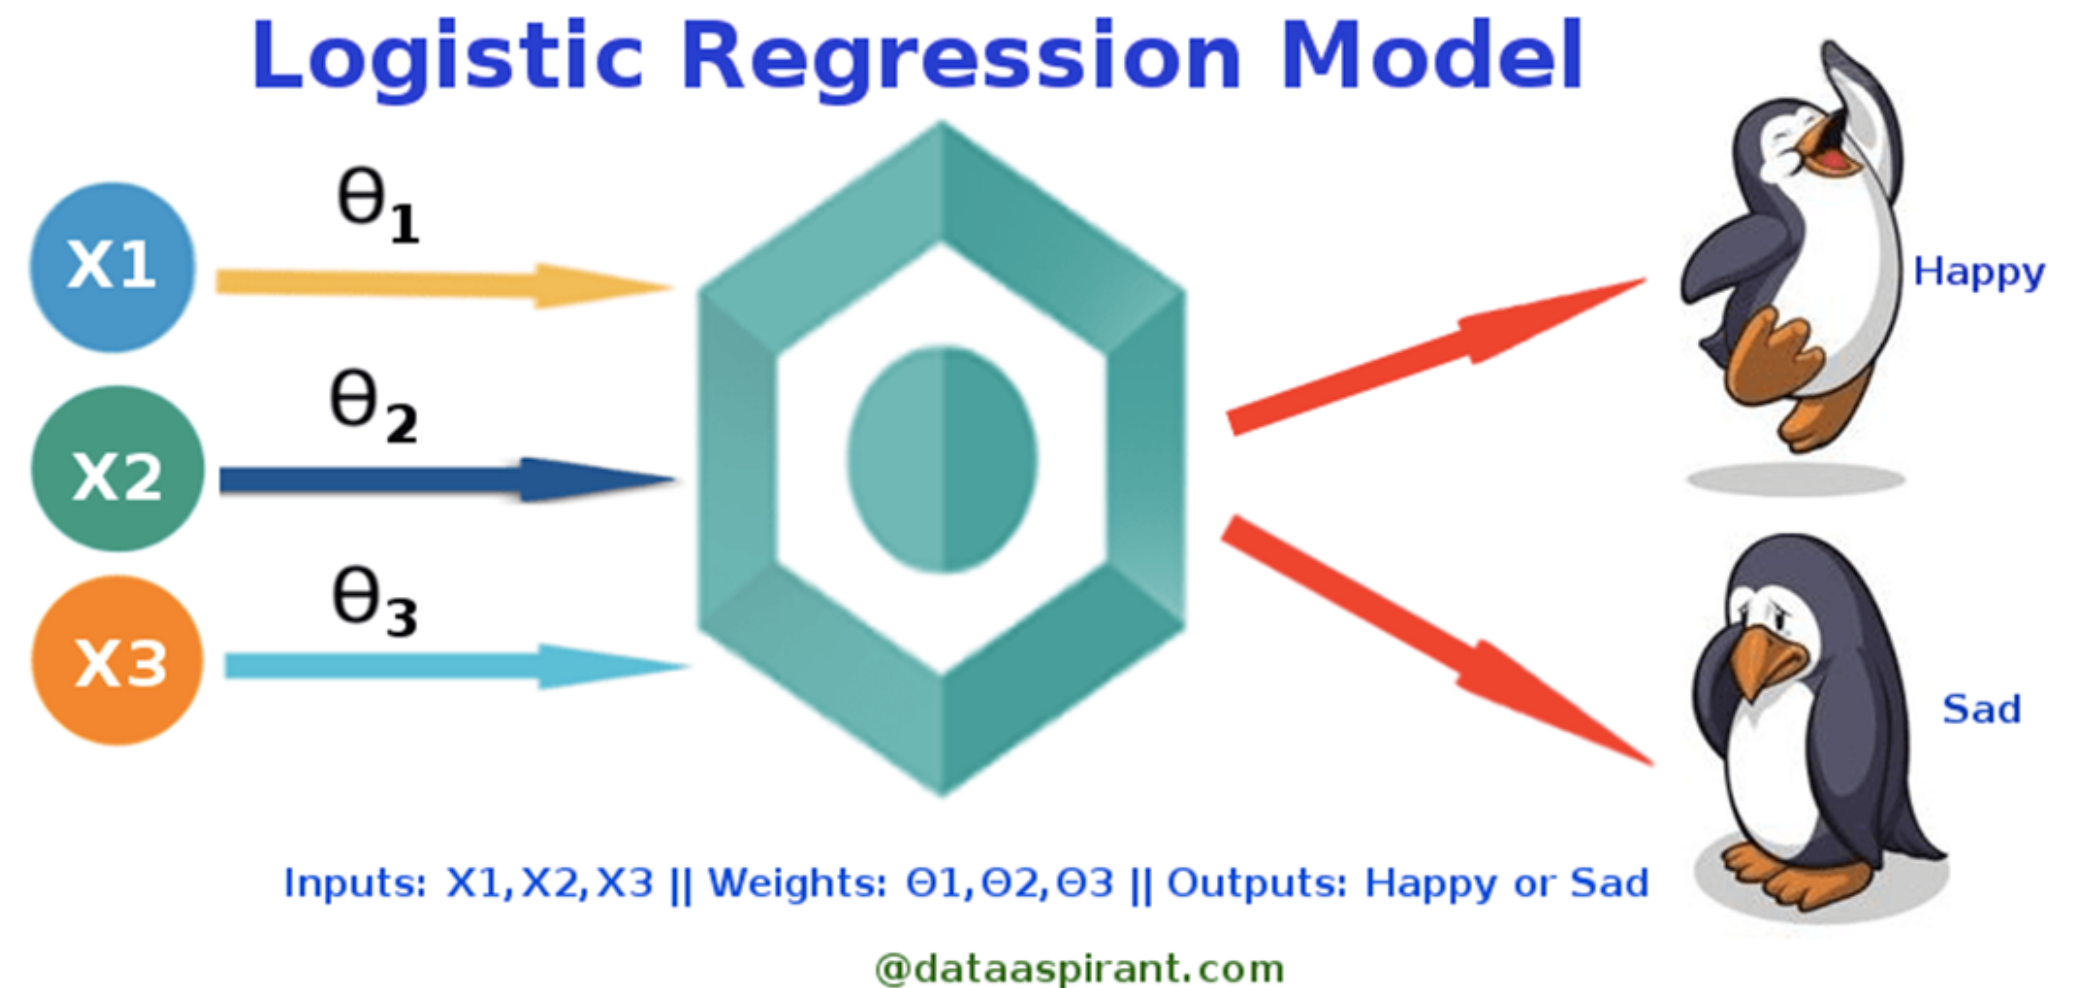
**Requires all data to be stored in memory to compute**  
**Can under perform with many variables**  
**Very sensitive to scale**

# LOGISTIC REGRESSION


## Binary Classification Algorithm

Takes in numerical inputs, takes a weighted sum of them and converts the result of that to a function that is "almost always" 0 or 1

**Hyperparameter:** Decision Boundary



# LOGISTIC REGRESSION



Studied	Slept	Passed
4.85	9.63	1
8.62	3.23	0
5.43	8.23	1
9.21	6.34	0

## Steps:

- Write usual linear combinations of input
- Pass result through sigmoid function
- Optimize parameters to minimize error
- Compare result to decision boundary
- Make Prediction

# LOGISTIC REGRESSION - LINEAR REGRESSION STEP


Studied	Slept	Passed
4.85	9.63	1
8.62	3.23	0
5.43	8.23	1
9.21	6.34	0

$$z = W_0 + W_1 \textit{Studied} + W_2 \textit{Slept}$$

**We do not optimize yet the parameters of this linear regression and we don't yet compare z to anything**

# LOGISTIC REGRESSION - SIGMOID FUNCTION

**This is where things become Nonlinear**



$$z = W_0 + W_1 \text{Studied} + W_2 \text{Slept}$$


$$f(z) = \frac{1}{1 + e^{-z}}$$

$$h(\vec{x}) = \frac{1}{1 + e^{-w_0 + w_1 x_1 + w_2 x_2}}$$

**This function goes from 0 to 1 and typically does so relatively "sharply"**  
**We now optimize  $W_1$  and  $W_2$  to "fit" the observations as well as possible**

# LOGISTIC REGRESSION - SIGMOID FUNCTION

This is where things become Nonlinear



$$z = W_0 + W_1 \text{Studied} + W_2 \text{Slept}$$

**For a set of the regression parameters, predictions are made and then compared with labels**

**Cost Function is used to evaluate "Penalty"**

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$


$$\text{Cost}(h_{\theta}(x), y) = -\log(h_{\theta}(x)) \quad \text{if } y = 1$$

$$\text{Cost}(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x)) \quad \text{if } y = 0$$

# LOGISTIC REGRESSION - SIGMOID FUNCTION

This is where things become Nonlinear

$$z = W_0 + W_1 \text{Studied} + W_2 \text{Slept}$$



$$f(z) = \frac{1}{1 + e^{-z}}$$

$$P(\text{class} = 1) = \frac{1}{1 + e^{-z}}$$




# LOGISTIC REGRESSION - SUMMARY



**Easy to implement**

**Trains quickly**


**Can be stacked to generate some really powerful models**



**Does not work as well for multi-class**

**Or at all for non-linearly separatable data**

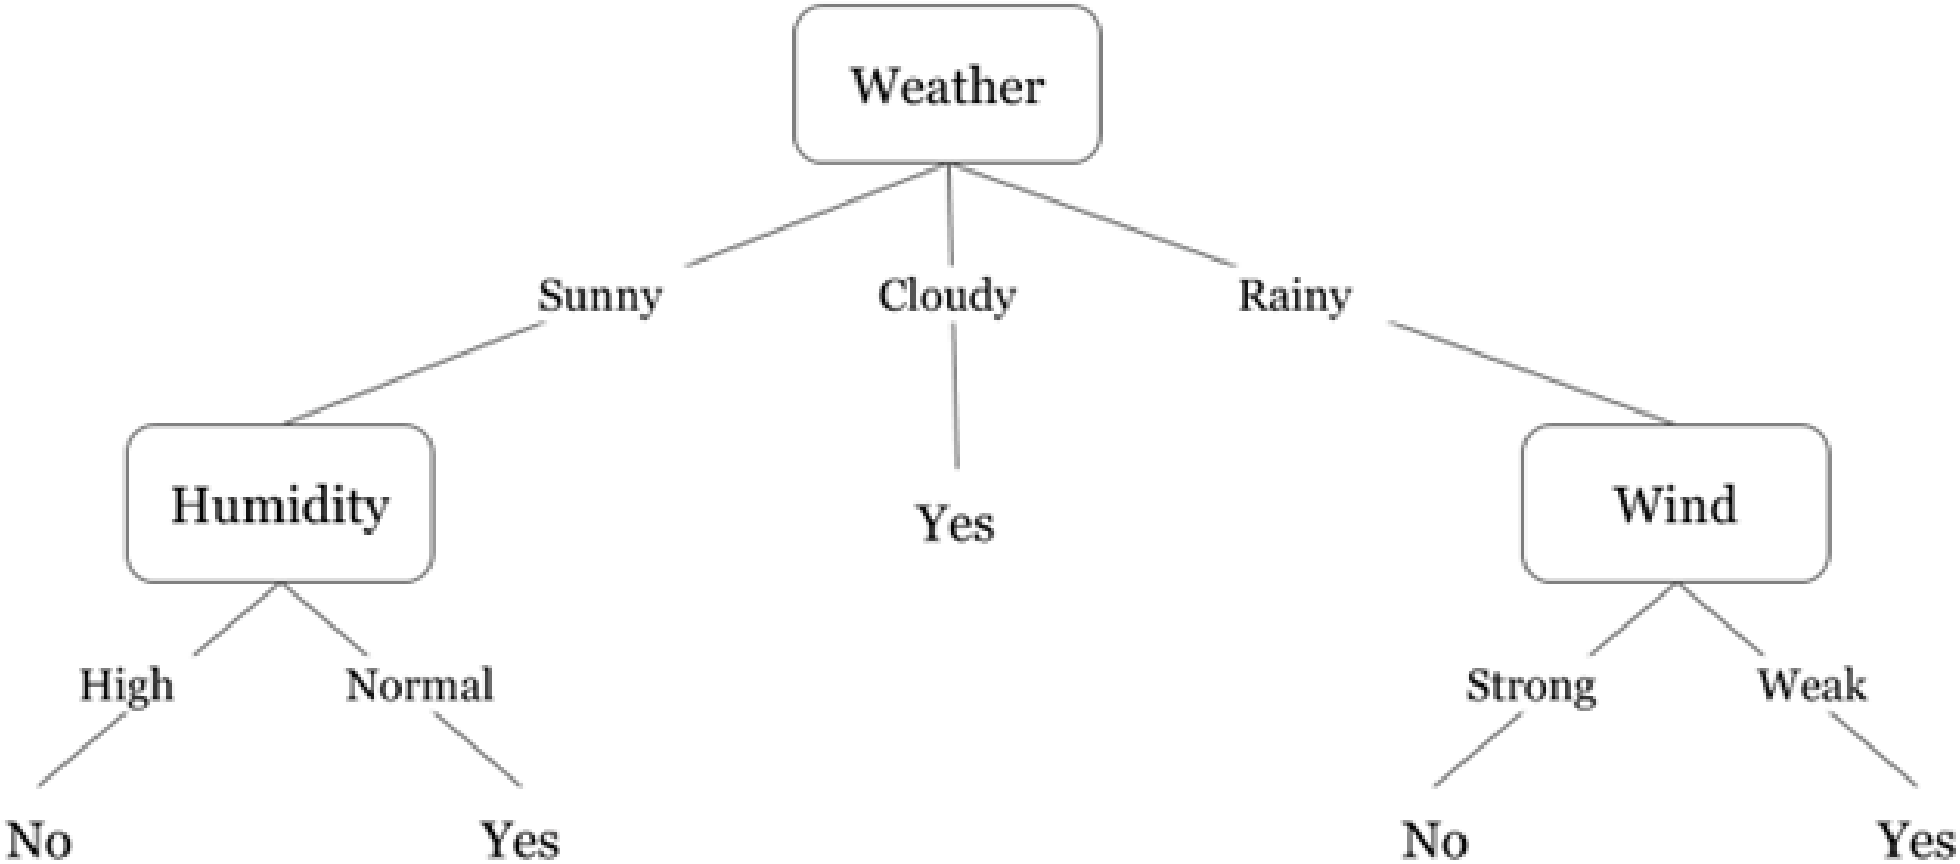
# DECISION TREES




- Each level of the tree asks a "question" about a specific feature and divides the paths into several nodes
- number of levels is known as the tree depth
- The question you ask each level is the one that gives you the most accurate answer if you were to stop then

# DECISION TREES

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY	PLAY GOLF
0	Rainy	Hot	High	False	No
1	Rainy	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Sunny	Mild	High	False	Yes
4	Sunny	Cool	Normal	False	Yes
5	Sunny	Cool	Normal	True	No
6	Overcast	Cool	Normal	True	Yes
7	Rainy	Mild	High	False	No
8	Rainy	Cool	Normal	False	Yes
9	Sunny	Mild	Normal	False	Yes
10	Rainy	Mild	Normal	True	Yes
11	Overcast	Mild	High	True	Yes
12	Overcast	Hot	Normal	False	Yes
13	Sunny	Mild	High	True	No




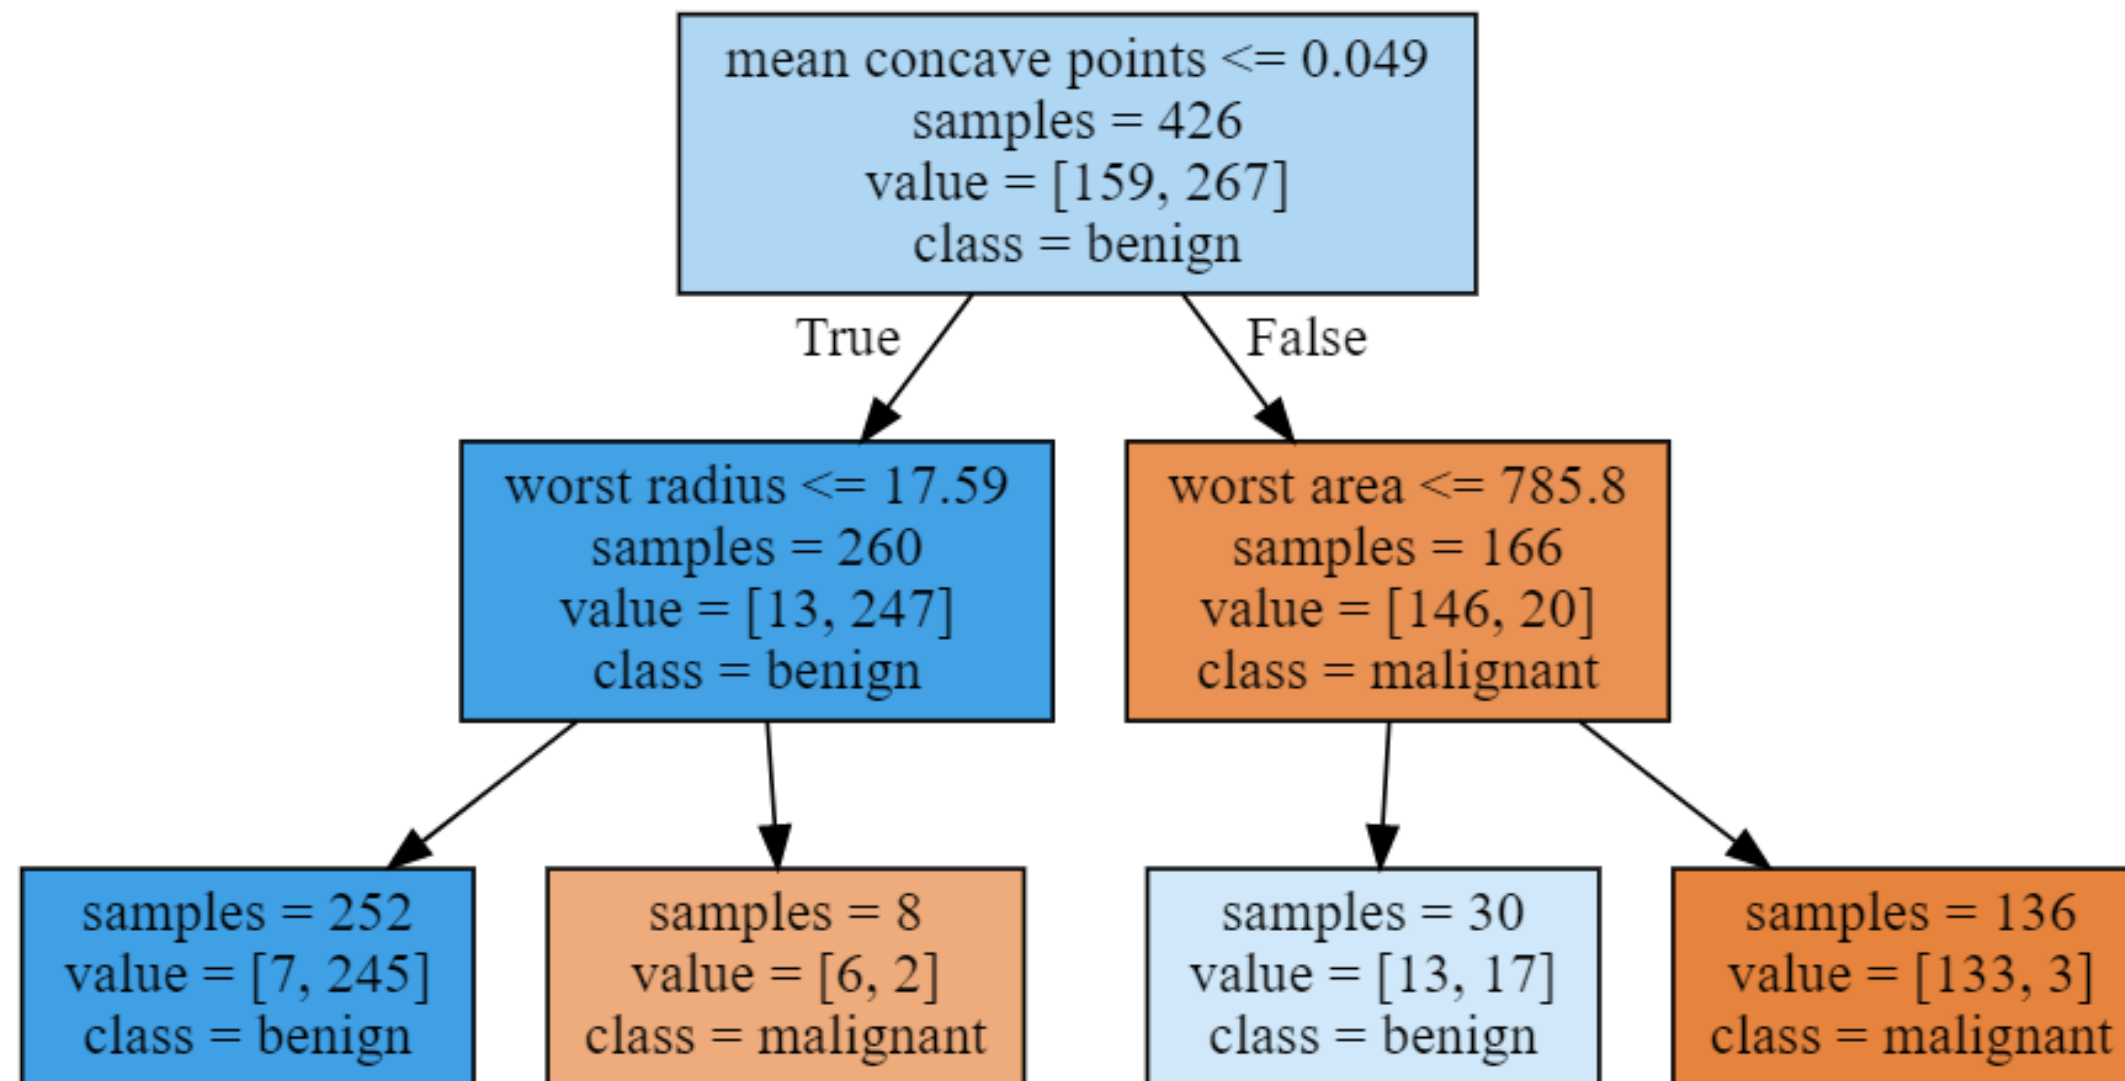
# DECISION TREES - SUMMARY



**Very robust to different types of features, including NaNs, categorical, numerical**  
**Handles non-linearity**

**Robust to outliers**

**Can be stacked into arguably the most useful models in modern ML**





**Very quickly overfit if hyperparameters are not controlled**


**Generalization Problems: what happens if feature outside possible ranges appears?**

**Online Learning problem: quickly change when re-trained**

# SUPPORT VECTOR MACHINE (SVM)



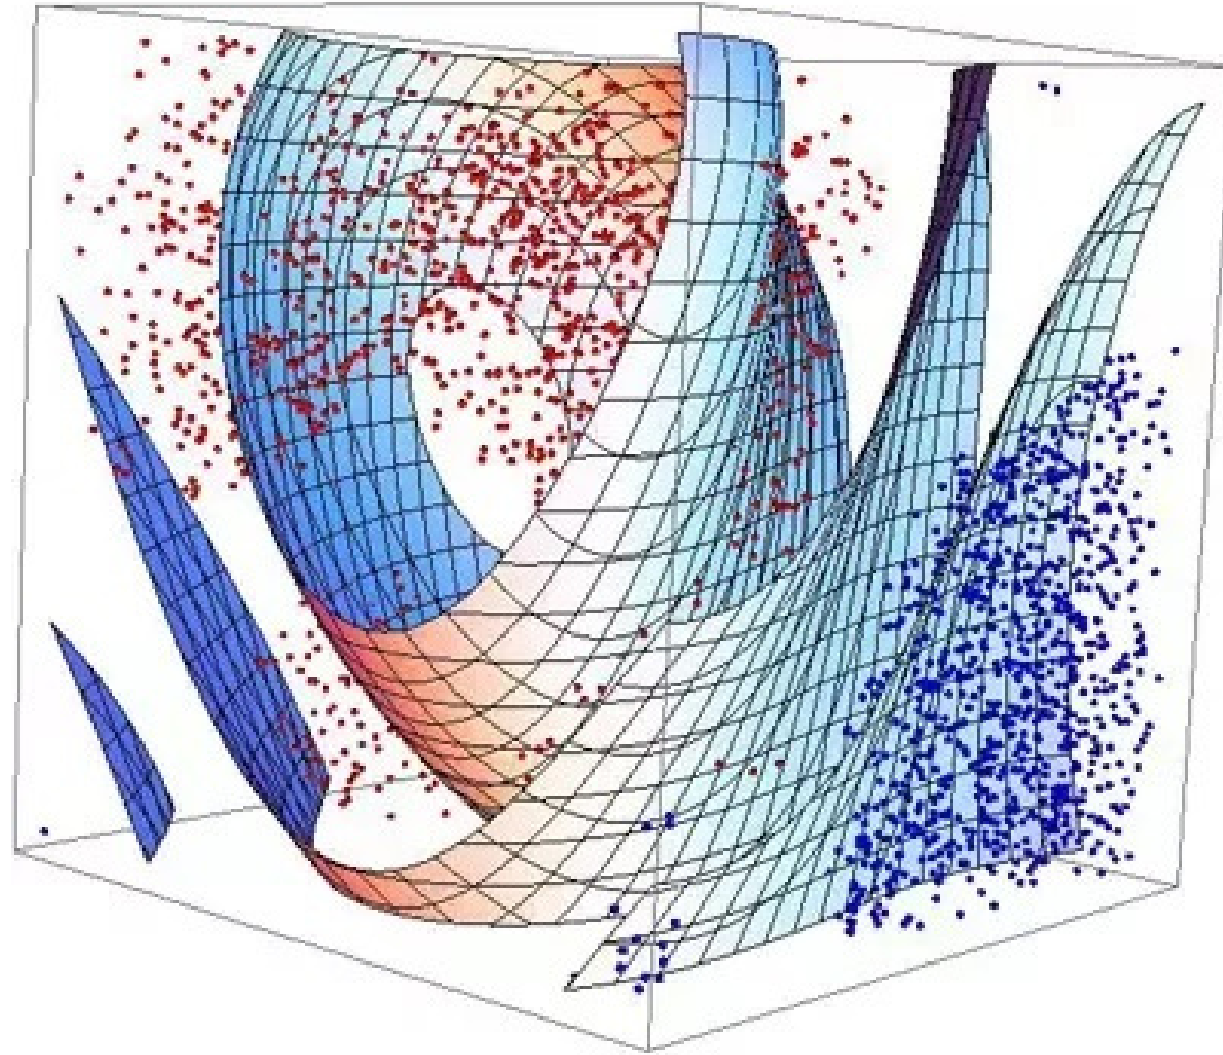
# SUPPORT VECTOR MACHINE (SVM)



**Our objective is to find a linear boundary that has the maximum margin, i.e the maximum distance between data points of both classes. Maximizing the margin distance provides some reinforcement so that future data points can be classified with more confidence.**

**We can force the SVM to maximize the distance to any point in any of the sets or add a loss function to allow but penalize some violations of the boundary**

# SUPPORT VECTOR MACHINE (SVM)



**What if there is no such boundary?**

**We can usually transform the underlying datapoints in such a way that a linear boundary exists. Then we revert that transformation and our boundary stops being linear, but is still effective for classification.**


**This is called a Kernel trick and usually requires very good understanding of the data**

# SVM - SUMMARY



**Can Model NonLinear Boundaries (if nonlinear version is used)**

**Unlikely to overfit**



**Memory intensive: is more successful on small dataset**



# CROSS VALIDATION

**Just to make the process more confusing!**

**Currently we've learnt that the train-test split is performed in a random way.**

**But are we leaving information on the table? How do we know a super representative point is being placed in the test dataset thus not allowing our algorithm to "learn"?**

**This is where Cross Validation comes in**

**Cros Validation is the process of performing **several different test-train splits****

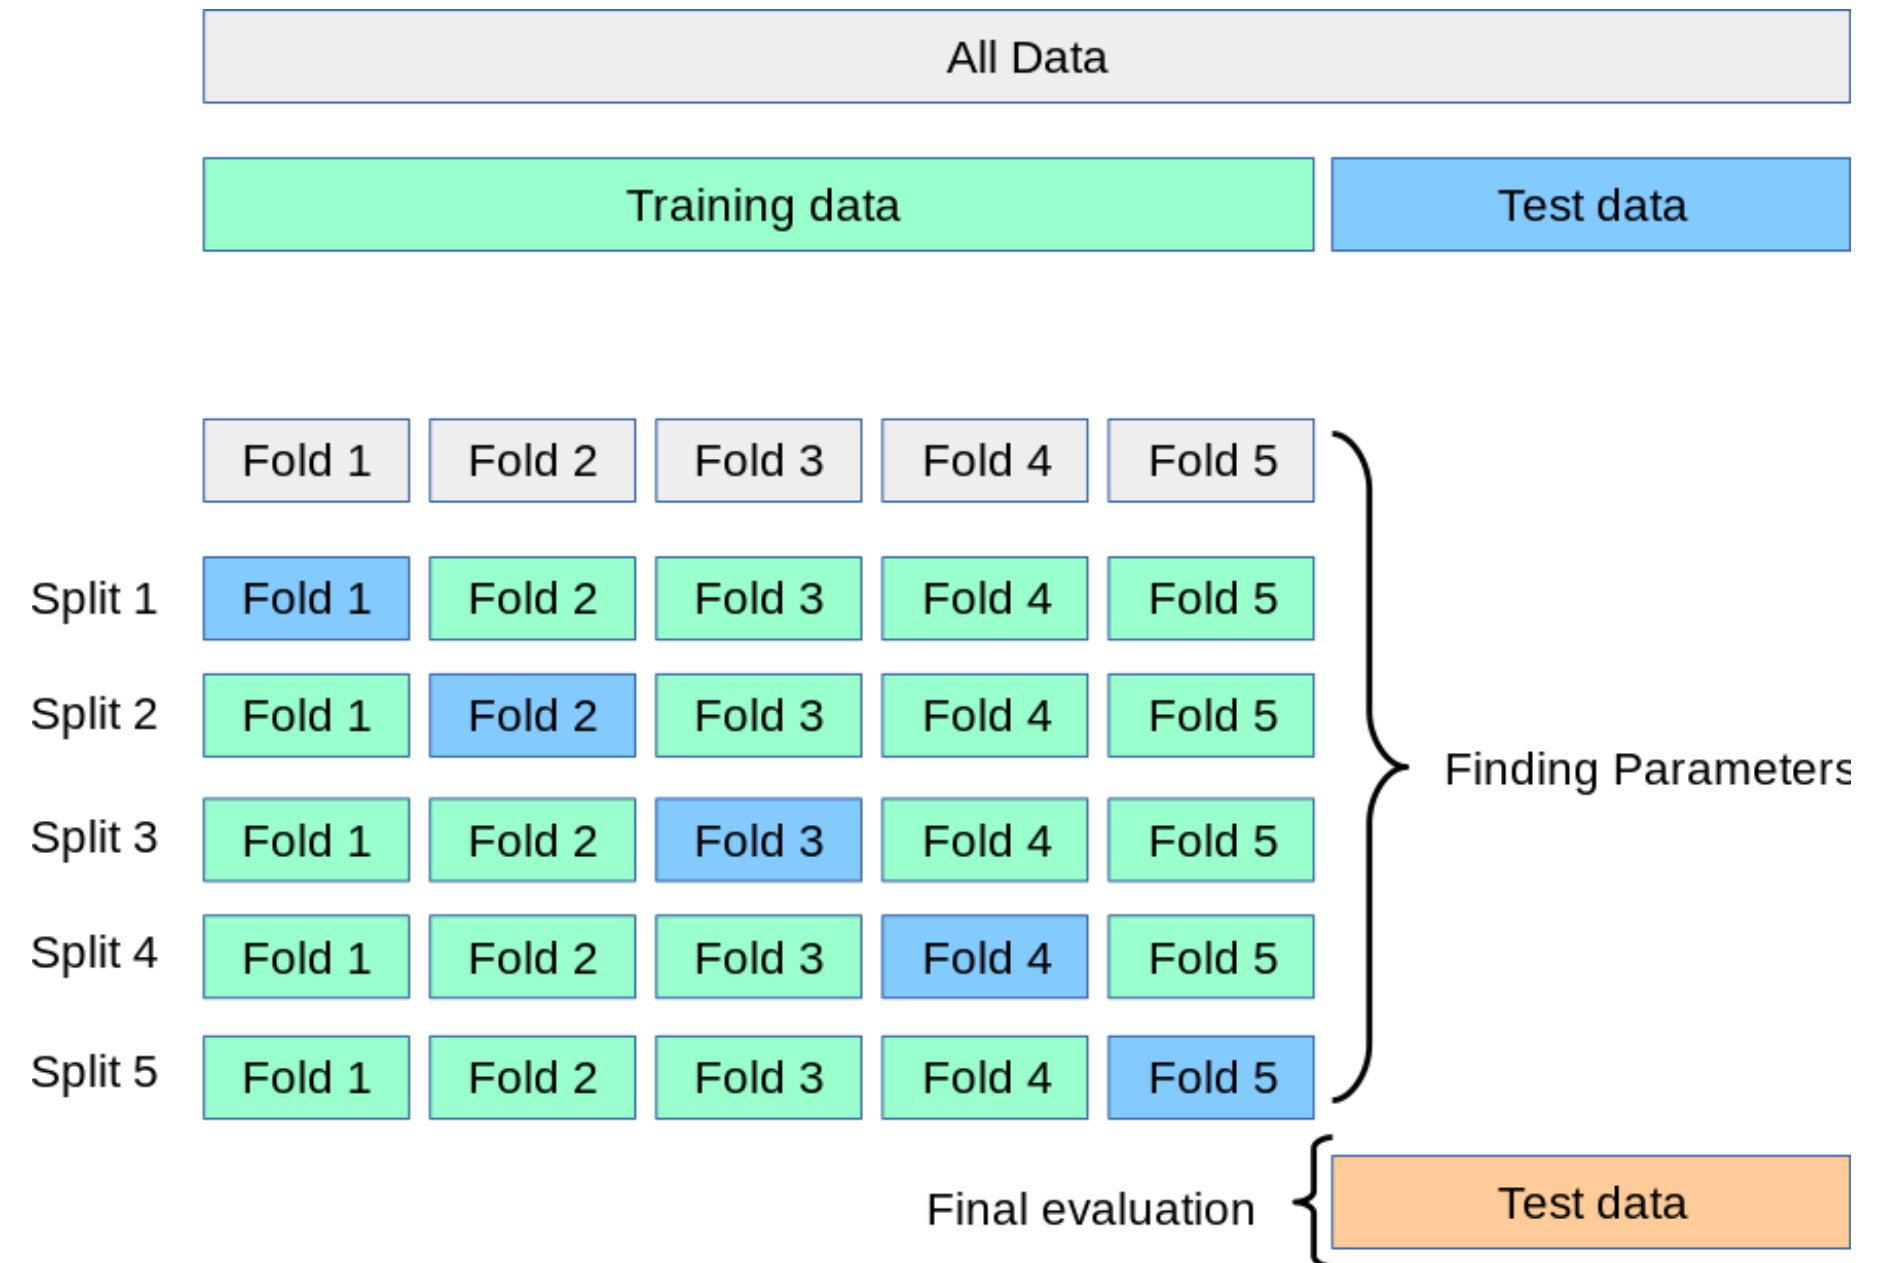
# CROSS VALIDATION

**Currently we've learnt that the train-test split is performed in a random way.**

**But are we leaving information on the table? How do we know a super representative point is being placed in the test dataset thus not allowing our algorithm to "learn"?**


**This is where Cross Validation comes in**

**Cros Validation is the process of performing several different test-train splits**



# CROSS VALIDATION

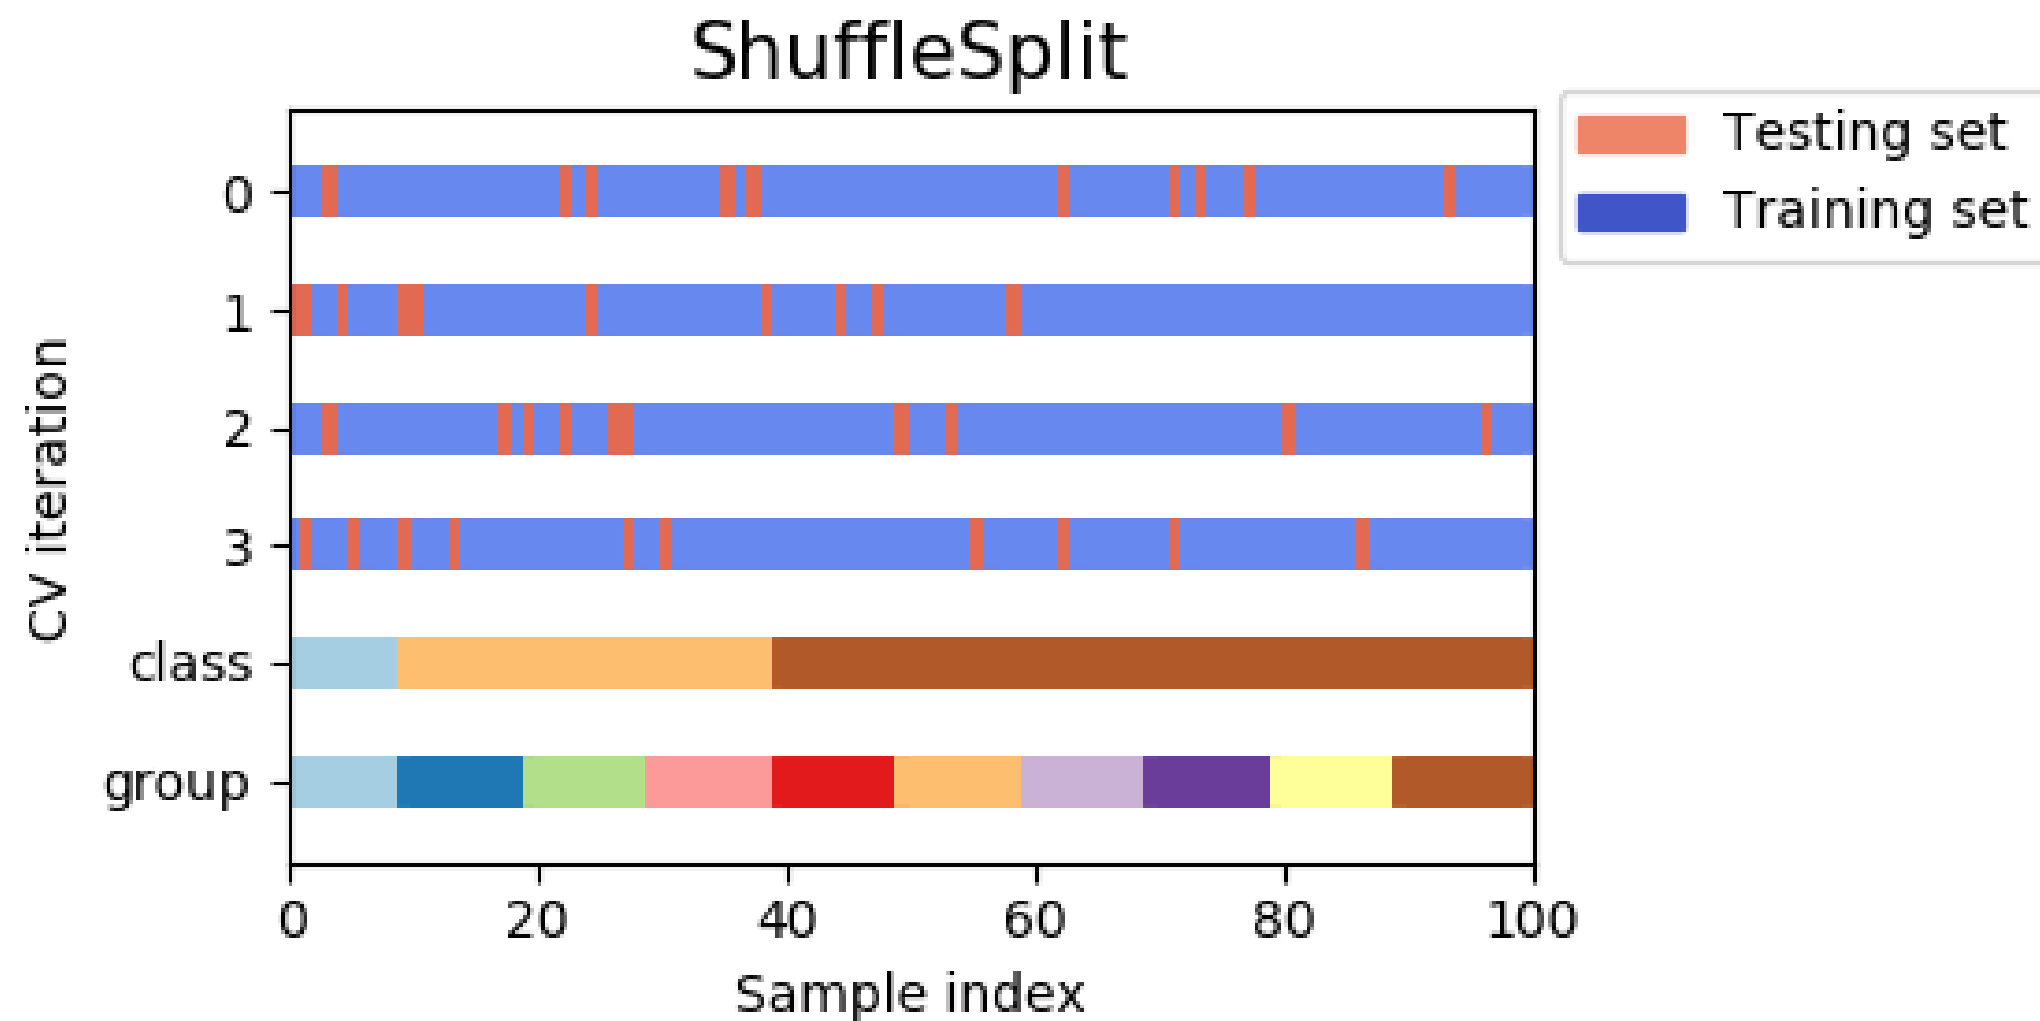
## K-Fold Cross Validation



**This is the simplest Cross Validation technique whereby data is split into k times into different Test-Train splits always preserving the length of each**

# CROSS VALIDATION


## Shuffle Split Cross Validation



**Performing a shuffled split further ensures more randomness is created when performing the split**

# CROSS VALIDATION

## Stratified KFold




**Whenever you have a severe class imbalance you may want to perform stratified sampling: you create your splits always preserving the same ratio of each class types (1/10)**

# CROSS VALIDATION


**How about the train-testing process for time series data?**

**Consider the problem if we used K-Fold**



# CROSS VALIDATION

**How about the train-testing process for time series data?**



**Super common mistake that I see in using Time Series Data is to use data from the future to build a model to predict the past!!**